

Device Embedding for Account Takeover Detection

A Hypothesis, Critique, Debate, and Empirical Resolution

Abstract

We investigate whether word-embedding models trained on per-account device login sequences can detect account takeover (ATO) by measuring the cosine distance between a candidate device's embedding and its target account's centroid. A proof-of-concept using FastText achieves ROC-AUC 0.836 on synthetic data. Systematic critique identifies ten theoretical and implementation weaknesses. A formal debate resolves six by argument; four require empirical evidence. A second experiment (Experiment 2) runs FastText, Word2Vec, and a binary OOV baseline under two corpus regimes (i.i.d. and Markov): FastText's character n-grams destroy per-account cluster structure (silhouette -0.051); Word2Vec produces near-perfect clusters (silhouette $+0.941$, AUC 0.982). A third experiment (Experiment 3) corrects two successive evaluation design flaws — all device IDs globally unique in Experiment 2 (making OOV detection equivalent to vocabulary membership lookup), and no legitimate new device enrollments in the negative class in the initial Experiment 3 run — and introduces feature-based embeddings (OS, browser, timezone, language, network type, screen resolution) in place of opaque device ID tokens. Under the corrected evaluation, the OOV binary baseline collapses to an analytically determined ceiling (AUC 0.750 on novel and spoof attacks) and is anti-correlated with fleet-reuse attack detection (AUC 0.250); these are not empirical results but the only possible outcomes for any binary OOV signal when 50% of the negative class consists of legitimate new device enrollment events. The OOV binary baseline is therefore abandoned as a risk scoring input. FastText trained on structured feature tokens (`feature_fasttext`) achieves AUC 0.985 on novel attacks and 0.920 on fleet-reuse attacks even with legitimate enrollment in the negative class. The recommended architecture is a two-path system: `feature_fasttext` embeddings for real-time profile-fit scoring; ID-based Word2Vec centroids (`id_w2v`) for offline fleet-reuse detection; and a per-account known-device set as an operational gate for step-up authentication.

1. Introduction

Account takeover fraud occurs when an attacker gains unauthorized access to a legitimate user's account, typically after credential theft. A primary behavioral signal is device novelty: genuine users log in from a consistent, recognizable set of devices, while attackers commonly arrive from devices that have no history with the account.

Simple novelty detection — “have we seen this device before?” — is a known baseline, but it is brittle. A returning attacker whose device was seen once months ago evades it. Richer signals require asking not just *whether* a device is new but *how well it fits the account's behavioral profile*.

The hypothesis explored here is that **word embeddings trained on per-account device sequences** can provide this richer signal. Each account’s login history is treated as a “sentence,” each device ID as a “token.” The distributional hypothesis — that tokens appearing in similar contexts acquire similar embeddings — predicts that devices belonging to the same account will cluster near each other in embedding space. A novel device that has never appeared in any account’s history will receive a generic embedding, landing far from the target account’s centroid and flagging the login as anomalous.

This report documents the complete arc of the investigation: initial implementation, systematic critique, formal debate, empirical resolution across two experiments, and a final architectural recommendation grounded in both experimental outcomes and production deployment constraints. The investigation extends beyond the initial debate/experiment cycle in two directions that materially changed the recommendation. First, production deployment analysis revealed that embedding spaces are not stable across training runs — each retraining event produces an incompatible coordinate frame, invalidating all stored account centroids simultaneously — and that the enrollment pressure from legitimate new device additions creates a continuous OOV false-positive window affecting millions of users at any given time. Second, the corrected evaluation design introduced in Experiment 3 places legitimate enrollment events in the negative class, exposing the OOV binary baseline as structurally incapable of distinguishing a new phone from an attacker at any threshold. These two constraints, taken together, drove the shift from OOV-binary to feature-profile-proximity as the primary real-time signal.

2. Experiment 1 — The Original PoC

2.1 Data generation

Three hundred synthetic accounts were generated, each with 3–10 distinct device IDs and a login history of 25–80 events. Device IDs are random alphanumeric strings (dev_ + 16 characters) with no semantic content. Login frequency follows a Zipf distribution — each account has a primary device used roughly 50% of the time and secondary devices appearing less frequently.

Takeover events were injected into 15% of accounts (45 accounts): each attacked account received one brand-new device ID that had never appeared anywhere in the training corpus. The evaluation set contains 300 legit test events (known devices, drawn from each account’s own fleet) and 45 attack events.


2.2 Model and training

A FastText skip-gram model (64 dimensions, window 5, 15 epochs, min_count=1) was trained on the per-account device histories as sentences. FastText was chosen specifically for its ability to embed out-of-vocabulary tokens via character n-gram decomposition — novel device IDs, never seen during training, would still receive a vector.


2.3 Scoring and evaluation

For each test event, the account’s training history was used to compute a centroid (frequency-weighted mean of its device embeddings). The cosine distance between the candidate device’s embedding and the centroid was the anomaly score. The Youden’s J optimal threshold was used to report precision and recall.

2.4 Results



Score distributions — original PoC



ROC curve — original PoC

Metric	Value
ROC-AUC	0.836
Known device dist. (mean $\pm$ std)	0.035 ± 0.027
Attack device dist. (mean $\pm$ std)	0.084 ± 0.051
Distance ratio (attack / known)	2.4x
Recall at Youden threshold	0.73
Precision at Youden threshold	0.35

The 2.4x mean distance gap and 0.836 AUC appeared to confirm the hypothesis. The UMAP visualization showed attack devices visually isolated from account clusters.

3. Critique

Ten weaknesses were identified, grouped by root cause. The four most consequential:

The stated mechanism may not be operating. The detection signal could be explained by a simpler phenomenon: novel attack devices are out-of-vocabulary (OOV), and OOV tokens receive generic embeddings that are generically anomalous — not specifically anomalous *for the target account*. A trivial vocabulary-membership lookup might achieve comparable AUC without embeddings. No such baseline was included in the original PoC.

Subword n-grams may harm rather than help. FastText decomposes every token into character n-grams. All device IDs share the `dev_` prefix, whose n-gram contributions are trained over the entire corpus — a cross-account average. The remaining 16 random characters produce n-gram overlaps by chance, not by account membership. Character n-grams might bleed cross-account signal into every token, collapsing exactly the per-account cluster structure the approach depends on.

The evaluation conflates two phenomena. Known devices are in-vocabulary and trained in their account’s co-occurrence context. Novel devices are OOV. Without comparing these two effects independently, the 0.836 AUC cannot be attributed to account clustering.

Synthetic data is too favorable. The i.i.d. device sampling assumption does not reflect real session autocorrelation. Real users use the same device across consecutive logins. Autocorrelation reduces the effective number of independent skip-gram training pairs, which may degrade embedding quality substantially.

Five additional issues were identified: Youden’s J threshold fitted on the test set (making precision/recall optimistic), ROC-AUC chosen over PR-AUC for a class-imbalanced problem, no silhouette validation of cluster structure in 64D, a misleading comment in the data generation code, and dead unreachable code in the UMAP visualization.

4. Defense

The defense accepted several critiques while rebutting others:

On the mechanism: The account centroid is account-specific; a binary OOV lookup is not. An attacker who reuses a device after its first login evades OOV detection permanently. Centroid distance continues to flag that device if it originated from a foreign account’s cluster. The two signals are complementary, not competing.

On subword n-grams: The shared `dev_` prefix is argued to be geometrically neutral — it appears identically in both device vectors and account centroids, canceling in the cosine dot product. The convergence of OOV devices to a consistent embedding region makes detection *easier* by concentrating the anomaly population, not harder.

On the evaluation: The production scenario requires scoring incoming devices against a centroid built from all known history. Known devices being in training is intentional, not leakage. The hypothesis claims “novel devices score higher,” not “account clustering alone causes this.”

On synthetic data: The i.i.d. simplification gives an upper-bound estimate. The direction of degradation under Markov sessions is predictable (AUC will drop), but whether it drops to uselessness requires measurement.

On ROC-AUC: The metric is correct for hypothesis validation — it answers “is there a discriminative signal?” before operational deployment concerns (class imbalance, cost structure) are relevant.

5. Debate

Ten points were argued to resolution. Four required empirical evidence; six were settled by argument.

Settled by argument:

- §3 (sparse co-occurrence): *Defense accepted*. Negative sampling actively pushes cross-account devices apart while the high-frequency pivot device pulls within-account devices together transitively. The mechanism is sounder than the critique assumed.
- §6 (Youden’s J): *Agreement*. ROC-AUC is unaffected; precision/recall at the Youden threshold should carry bootstrap confidence intervals or be omitted.
- §7 (ROC-AUC vs PR-AUC): *Agreement*. ROC-AUC is correct for phase-one hypothesis validation. PR-AUC is the production deployment metric.
- §4 (centroid norm correction): *Both concede*. The proposed formula $\text{score} / (||\mu|| + \epsilon)$ has the wrong sign for its stated intent. The correct approach is to output score and $||\mu||$ as separate quantities; the result should settle empirically whether $||\mu||$ predicts false positives.
- §9 (UMAP): *Agreement*. Silhouette score in 64D is the correct cluster validation.
- §10 (incremental update): *Defense concedes scope*. The signal’s FPR floor is bounded below by legitimate device turnover; this is a primary operational constraint that belongs in the README, not a footnote.

Requiring empirical resolution:

- §1 (mechanism): Score cross-account in-vocab devices against a target account’s centroid. If cross-account $\approx$ OOV attack, account clusters are real. If cross-account $\approx$ known, OOV is the dominant signal.
- §2 (subword n-grams): Run the experiment with Word2Vec (no subword n-grams) to isolate the character n-gram effect on AUC.
- §5 (OOV baseline): Add a binary OOV baseline to the evaluation. If FastText’s AUC significantly exceeds it, embeddings are adding value.
- §8 (Markov sessions): Implement a Markov session corpus and compare AUC. If FastText degrades substantially while simpler methods hold, i.i.d. is confirmed as an upper bound.

6. Experiment 2 — Empirical Resolution

6.1 Design

Three models were evaluated:

Model	OOV handling	Hypothesis tested
FastText (skip-gram, dim=64)	Subword n-gram averaging	Original PoC approach
Word2Vec (skip-gram, dim=64)	Global mean vector fallback	Embedding without subword noise
OOV binary baseline	1.0 if unseen, 0.0 if in-vocab	Trivial novelty detection

Four event classes were scored for each model:

- **Known:** a device from the account’s own training history
- **Cross-account (in-vocab):** a device from a *different* account’s training history
- **OOV attack:** a brand-new device ID never seen in any account’s history

Two corpus modes were used:

- **i.i.d.:** Zipf-weighted independent draws (original PoC assumption)
- **Markov:** sticky device-switching with stay probability 0.70, creating session runs

Bootstrap confidence intervals (N=1,000, 95%) were computed for all AUC estimates. Silhouette score was computed in 64D using cosine distance over account-labeled device embeddings. The centroid norm $||\mu||$ was tested for association with false positives via Mann-Whitney U.

6.2 The mechanism test (Point 1)

The single most important test: does account clustering exist, and which class is actually driving the detection signal?

Fig 1 — Mechanism Test: Account Clustering vs OOV Detection (Point 1)




Fig 1 — Three-class mechanism test

Event class	Mean cosine dist (i.i.d.)	Mean cosine dist (Markov)
Known (own account)	0.035	0.098
Novel OOV attack	0.097	0.182
Cross-account (in-vocab)	0.131	0.241

Cross-account in-vocab devices score higher than OOV attack devices. The AUC for “cross-account vs. OOV attack” is **0.375** (i.i.d.) and **0.279** (Markov) — both well below 0.5, confirming the ranking is inverted.

This result resolves Point 1 in an unexpected direction. Account clustering *is* real — the embedding space does separate accounts, as demonstrated by cross-account devices scoring higher than own-account devices. But the OOV mechanism works in the *opposite* direction from what the critique predicted: novel attack devices do not land far from all centroids. FastText’s subword averaging causes them to converge toward the global mean of the embedding space, which happens to be angularly *close* to trained account centroids. The detection signal comes from known devices being pulled *especially close* to their centroid by skip-gram training, not from attack devices being pushed far away.

6.3 Model comparison (Points 2, 5, 7)




Fig 2 — Model AUC comparison with bootstrap CIs

Model	AUC i.i.d. [95% CI]	AUC Markov [95% CI]
FastText centroid	0.910 [0.848–0.963]	0.832 [0.747–0.916]
Word2Vec centroid	0.982 [0.940–1.000]	0.962 [0.906–1.000]
OOV binary baseline	0.989 [0.962–1.000]	0.956 [0.907–0.990]

Word2Vec outperforms FastText by 7.2 AUC points on the i.i.d. corpus. The OOV binary baseline outperforms FastText by 7.9 points. FastText, the model chosen specifically for its OOV embedding capability, produces the weakest results of the three.

6.4 Cluster structure in 64D (Point 9)

The silhouette score is the direct test of whether per-account clusters exist in the full embedding space — the foundational claim of the hypothesis.




Fig 3 — UMAP and silhouette scores for FastText vs Word2Vec

Model	64D Silhouette	Interpretation
FastText	-0.051	No account cluster structure
Word2Vec	+0.941	Near-perfect account separation

FastText’s silhouette of -0.051 means that on average, devices are *closer* to devices from other accounts than to devices from their own account. The cross-account signal in the embedding space is negative. FastText’s subword n-grams have collapsed the block-diagonal co-occurrence structure that skip-gram negative sampling creates, by forcing all device vectors toward a shared character-n-gram attractor regardless of account membership.

Word2Vec’s silhouette of $+0.941$ is near the theoretical maximum. Without subword interference, skip-gram negative sampling produces tight, well-separated per-account clusters. This is the empirical validation the hypothesis required — and it is achieved by the model *without* the subword mechanism.

6.5 Session autocorrelation (Point 8)

Fig 4 — i.i.d. vs Markov: Session Autocorrelation Degrades FastText Most (Point 8)




Fig 4 — i.i.d. vs Markov degradation

Model	i.i.d.	Markov	Δ
FastText	0.910	0.832	-0.078
Word2Vec	0.982	0.962	-0.020
OOV binary	0.989	0.956	-0.033

FastText degrades four times faster than Word2Vec under Markov session structure. Session autocorrelation reduces the effective number of distinct skip-gram training pairs per account. FastText is doubly harmed: fewer informative co-occurrence pairs *and* subword noise. Word2Vec’s tight cluster structure (silhouette 0.927 under Markov, vs 0.941 under i.i.d.) is nearly unaffected by the reduced pair count. The i.i.d. AUC of 0.836 in the original PoC was an optimistic upper bound on what FastText achieves under realistic session patterns.

6.6 Signal stability across account history lengths (Point 3)

Fig 5 — Does Signal Survive Thin Histories? Stratified AUC (Point 3)




Fig 5 — Stratified AUC by history length

FastText’s AUC is surprisingly stable across thin and rich accounts on the i.i.d. corpus (range 0.870–0.950), partially vindicating the defense’s position on co-occurrence sparsity. Word2Vec achieves ≥ 0.940 across all strata under i.i.d. Both models degrade more on thin accounts under Markov, as expected. The OOV binary baseline is notably less stable under Markov for thin accounts: session runs limit vocabulary coverage, meaning some device IDs that should be recognized as in-vocab appear only once in training and are intermittently misclassified as OOV.

This identifies the one practical advantage of the centroid-based approach over the binary baseline: **for thin histories under realistic session patterns, Word2Vec centroid distance is more reliable than the OOV lookup.**

6.7 Centroid norm as a confidence signal (Point 4)

Fig 6 — Centroid Norm vs False Positives: $||\mu||$ is not predictive (Point 4)




Fig 6 — Centroid norm vs false positives

Corpus	Mann-Whitney p (	
i.i.d.	0.117	Not significant
Markov	0.573	Not significant

The centroid norm $||\mu||$ does not differ significantly between accounts that produce false positives and accounts that do not. The critique’s proposed formula score $||\mu|| / (||\mu|| + \epsilon)$ was shown during the debate to have the wrong sign; this empirical result confirms that $||\mu||$ is not a useful confidence feature for this data at all. Account history length is a better (if imperfect) proxy for centroid reliability, as evidenced by the stratified AUC results.

7. Experiment 3 — Corrected Evaluation and Feature-Based Embeddings

7.1 Two evaluation design flaws corrected

Before reporting Experiment 3 results, both evaluation design flaws that preceded them must be stated explicitly, because the corrections are load-bearing for interpreting every AUC number in this section.

Flaw 1 — Experiment 2: The synthetic data assigned every account a unique set of device IDs. There were no cross-account devices. The global OOV binary baseline achieved AUC 0.989, but this was entirely explained by the data design: all attack devices had novel IDs that had never appeared anywhere in the training corpus, and there were no legitimate new device enrollments in the evaluation negative class. Detecting an attack was equivalent to detecting a globally unseen device ID — a trivially separable task that does not resemble the real threat model. The 0.989 AUC is an artifact of this design, not a measurement of the signal’s discriminative power in production.

Flaw 2 — initial Experiment 3 run: Correcting Flaw 1 introduced cross-account fleet devices into the training vocabulary, but the evaluation negative class still consisted entirely of returning known-device sessions. There were no legitimate new device enrollments in the evaluation. Any signal that scores all OOV devices as suspicious achieved AUC 1.000 by construction, because there were no OOV legitimate events to confuse it. This is the exact scenario production systems face continuously: users buy new phones, switch browsers, and travel across timezones — and a signal that flags every new device as an attack is operationally useless.

The corrected design (Experiment 3, final run): The negative class contains two event types at equal weight:

- *Returning*: a known device ID from the account’s training history, with the account’s primary feature profile. Label = 0.
- *Enrollment*: a new device ID (globally OOV, not in any training vocabulary), with a feature profile consistent with the account’s primary profile — same OS, browser, timezone, and language; randomized network type and screen resolution. Label = 0.

This correction changes the story fundamentally. Any binary signal that fires on all OOV devices now faces a hard ceiling: it cannot distinguish a legitimate enrollment from an attack at any threshold. Only signals that measure something about the device’s feature profile relative to the account’s behavioral history can escape that ceiling.

7.2 Data redesign: fleet corpus and feature representation

The Experiment 3 corpus was built around two structural additions absent from Experiments 1 and 2.

Fleet corpus: 400 accounts were generated, each with a primary device profile defined across six dimensions (OS, browser, timezone, language, network type, screen resolution) and 2–4 known devices assigned at setup. Eighty shared fraud fleet devices — with attacker profiles (Windows or Linux OS, UTC+5 or UTC+8 timezone, non-English language) — were injected into 25% of accounts during training. Fleet device IDs therefore appear in the global Word2Vec training vocabulary through cross-account co-occurrence, but they are not in any individual target account’s known-device set. This models the real threat: an attacker’s device has been seen somewhere in the system before, but it has no legitimate history with the specific account being targeted.

Feature token representation: Rather than treating device IDs as opaque strings, each login event is mapped to a sequence of six structured feature tokens. A device running iOS Safari in UTC-5 with English language settings produces tokens such as `os_ios`, `browser_safari`, `tz_utc-5`, `lang_en`, `net_wifi`, `screen_1920x1080`. The complete feature vocabulary across all six dimensions contains approximately 30 distinct tokens (OS: ~4 values, browser: ~5, timezone: ~8, language: ~6, network type: ~4, screen resolution: ~6). This bounded vocabulary is the key property that distinguishes feature-based embeddings from ID-based embeddings: a new device is not out-of-vocabulary because it has features, not an opaque identifier.

7.3 Three attack types

Three attack types were constructed to cover distinct threat models:

novel: New device ID (globally OOV, never seen in any training corpus) combined with a clearly foreign feature profile — different OS family, far timezone (UTC+5 or UTC+8 when the account is UTC-5), non-English language. This is the clearest attack signal: both the device ID and its behavioral fingerprint are foreign to the account.

fleet: An existing fleet device ID — one of the 80 attacker-profile devices present in the global Word2Vec training vocabulary from cross-account injection, but not in the target account’s known-device set — combined with the attacker’s profile. This models device reuse across victims: the device is globally known but account-foreign. The binary OOV baseline returns 0.0 (device is in vocabulary), so the baseline cannot detect this attack type at all.

spoof: New device ID (globally OOV) combined with a victim-mimicking feature profile — same OS, browser, and language as the account’s primary profile; different timezone only. This models a sophisticated attacker who has researched the victim’s device characteristics and attempts to blend in. Only signals that measure fine-grained feature proximity can detect the timezone mismatch.

7.4 Six signals evaluated

Six signals were evaluated across all three attack types and both corpus modes:

global_oov: Binary signal returning 1.0 if the device ID is not present in the Word2Vec training vocabulary (i.e., the global model has never seen this device across any account). Returns 0.0 for fleet devices, which are in the vocabulary. Returns 1.0 for enrollment events, which are also OOV. Cannot distinguish enrollment from novel attack at any threshold.

account_oov: Binary signal returning 1.0 if the device ID is not in the per-account known-device set. Unlike `global_oov`, this fires on fleet devices (fleet devices are not in the target account’s known-device set). But it also fires on enrollment events (new device IDs are not in the known-device set). Cannot distinguish enrollment from any attack type at any threshold.

id_w2v: Cosine distance from the device’s ID-based Word2Vec embedding to the account’s ID-based Word2Vec centroid. OOV devices (novel attacks, spoof attacks, enrollment events) all receive the global mean vector; their cosine distance to the account centroid is therefore identical across all three event types. Fleet devices receive genuine embeddings positioned by cross-account co-occurrence — typically far from any single account’s centroid — making this the fleet-detection signal.

feature_w2v: Cosine distance from the device’s feature-based Word2Vec embedding to the account’s feature-based Word2Vec centroid. Each device’s embedding is the mean of its six feature token embeddings (OS, browser, timezone, language, network type, screen resolution). Because the feature vocabulary is bounded (~30 tokens), all devices — including new ones — receive positioned embeddings based on their feature content, not their device ID. An enrollment device with a profile matching the account’s primary profile receives an embedding close to the account centroid; an attack device with a foreign profile lands far.

feature_fasttext: Cosine distance to the account’s feature-based FastText centroid. Identical in structure to `feature_w2v` but trained with FastText on the same structured feature tokens. FastText’s character n-gram mechanism operates on meaningful token prefixes (`os_`, `browser_`, `tz_`) rather than random character sequences, reinforcing feature-dimension clustering. Critically, FastText embeds any feature token at inference time via n-gram averaging — including tokens not seen during training (new OS versions, browser families, timezone variants) — without requiring an OOV fallback to the global mean vector. This is the key production advantage over `feature_w2v`.

feature_novelty: Binary signal returning 1.0 if the device’s exact six-dimension feature-profile tuple has not been observed anywhere in the account’s training history. Falls between the pure OOV signals and the embedding-based signals: it measures profile novelty rather than ID novelty, so it does not collapse to the 0.750 OOV ceiling, but it is sensitive to granularity — enrollment devices with randomized network type and screen resolution often produce novel tuples even though their core profile (OS, browser, timezone, language) matches the account.

7.5 Results

i.i.d. corpus:

Signal	novel	fleet	spooF
global_oov	0.750 [0.711–0.787]	0.250 [0.211–0.287]	0.750 [0.711–0.787]
account_oov	0.750 [0.711–0.787]	0.750 [0.711–0.787]	0.750 [0.711–0.787]
id_w2v	0.691 [0.627–0.758]	0.891 [0.843–0.929]	0.691 [0.627–0.758]
feature_w2v	0.984 [0.971–0.994]	0.914 [0.880–0.945]	0.817 [0.765–0.868]
feature_fasttext	0.985 [0.972–0.994]	0.920 [0.885–0.949]	0.798 [0.747–0.851]
feature_novelty	0.800 [0.760–0.838]	0.763 [0.712–0.813]	0.800 [0.760–0.838]

Markov corpus:

Signal	novel	fleet	spooF
global_oov	0.750 [0.711–0.787]	0.250 [0.211–0.287]	0.750 [0.711–0.787]
account_oov	0.750 [0.711–0.787]	0.750 [0.711–0.787]	0.750 [0.711–0.787]
id_w2v	0.696 [0.624–0.762]	0.856 [0.800–0.906]	0.696 [0.624–0.762]

Signal	novel	fleet	spooF
feature_w2v	0.993 [0.984–0.999]	0.913 [0.873–0.945]	0.809 [0.757–0.856]
feature_fasttext	0.988 [0.976–0.997]	0.901 [0.858–0.935]	0.785 [0.729–0.835]
feature_novelty	0.791 [0.750–0.825]	0.766 [0.718–0.807]	0.791 [0.750–0.825]




Fig E3-1 — AUC heatmap by signal and attack type, i.i.d. and Markov




Fig E3-3 — Signal comparison by attack type with 95% CI, i.i.d. corpus

7.6 Key findings from Experiment 3

Finding 1 — The 0.750 ceiling is analytical, not empirical.

Any binary signal that fires on all OOV events produces AUC 0.750 by construction when 50% of the negative class is enrollment events. The derivation is direct. With the corrected evaluation design, all attack events score 1.0 under `account_oov` (attacks are not in the known-device set) and all enrollment events also score 1.0 (new device IDs are not in the known-device set). Returning known-device events score 0.0. The ROC curve has three points: (0, 0) at threshold above 1.0; (0.5, 1.0) at threshold = 1.0, which flags all OOV events — recovering all attacks (TPR = 1.0) but also flagging all enrollment events (FPR = 0.5, since enrollment is half the negative class); and (1.0, 1.0) at threshold = 0.0. The area under this piecewise linear curve is $0.5 \times 0.5 \times 1.0 + 0.5 \times 1.0 = 0.75$. This is not an empirical result. It is the only possible result for any binary OOV signal under this evaluation design, regardless of the signal's implementation. The same derivation applies to `global_oov` on novel and spoof attacks, since enrollment events are also globally OOV.

Finding 2 — global_oov on fleet is inverted (AUC 0.250).

Fleet devices are present in the Word2Vec training vocabulary — they were injected into training sessions across 25% of accounts — so global_oov returns 0.0 for fleet attack events. Enrollment events are globally OOV (new device IDs), so global_oov returns 1.0 for enrollment events. The ROC curve for fleet detection therefore has the following structure: at threshold = 1.0, the signal flags all enrollment events (none of which are fleet attacks) and zero fleet attacks — TPR = 0.0, FPR = 0.5. At threshold = 0.0, all events are flagged — TPR = 1.0, FPR = 1.0. AUC = 0.25. The signal is not merely uninformative about fleet attacks; it actively ranks enrollment events as more suspicious than actual fleet attacks. This is the key failure mode the original hypothesis was designed to address: a pure vocabulary membership check cannot distinguish a device that is globally known but account-foreign from a device that is globally new but account-consistent.

Finding 3 — feature_w2v and feature_fasttext are the only signals that structurally separate enrollment from attack.

Enrollment devices have feature profiles matching the account's primary profile — same OS, browser, timezone, and language. Their six feature token embeddings are drawn from high-frequency tokens in the account's training corpus, and the mean of those six embeddings lands close to the account's feature centroid. Novel attack devices have foreign profiles drawn from the attacker distribution (different OS, far timezone, non-English language); their feature embeddings land far from the account centroid. The result is AUC 0.984 (i.i.d.) / 0.993 (Markov) for feature_w2v and AUC 0.985 (i.i.d.) / 0.988 (Markov) for feature_fasttext on novel attacks, even with enrollment in the negative class. No other signal achieves this separation: global_oov and account_oov collapse to the 0.750 ceiling because they cannot see feature content; id_w2v collapses to ~0.691–0.696 on novel/spoof because OOV attack and OOV enrollment devices both receive the global mean ID vector, making them indistinguishable; feature_novelty degrades to ~0.800 because enrollment events with randomized network type and screen resolution often produce novel tuples.

Finding 4 — id_w2v is the fleet/reuse detection signal.

Fleet devices have ID-based Word2Vec embeddings that were positioned through co-occurrence with multiple different accounts during training. A device that appeared in sessions for accounts A, B, and C — each with different primary devices — is embedded in a region of the space that is not close to any single account's centroid. When it appears in account D's evaluation, the cosine distance to account D's ID centroid is high. AUC 0.891 (i.i.d.) / 0.856 (Markov) on fleet attacks. id_w2v cannot distinguish novel or spoof attacks from enrollment events: all three produce OOV device IDs, all three receive the global mean vector, and the global mean's cosine distance to the account centroid is nearly identical across these event types (AUC ~0.691–0.696, slightly below the analytical 0.750 ceiling because the global mean's position is not perfectly centered). id_w2v is a fleet-detection signal, not a general ATO signal.

Finding 5 — feature_fasttext matches feature_w2v within confidence intervals on all attack types; the meaningful production differentiator is OOV token handling.

The AUC differences between feature_fasttext and feature_w2v are not statistically distinguishable at 95% confidence on novel (0.985 vs 0.984 i.i.d.), fleet (0.920 vs 0.914 i.i.d.), or spoof (0.798 vs 0.817 i.i.d.) attacks. On spoof attacks, feature_fasttext is marginally weaker than feature_w2v: a spoofed device presents the same OS, browser, and language tokens as the victim account — these tokens share os_, browser_, and lang_ n-grams with legitimate feature tokens in the account's training history. FastText's n-gram averaging pulls the spoofed profile's embedding slightly closer to the account centroid than Word2Vec would, because Word2Vec treats each token as an atomic unit. The meaningful production distinction is what happens when a feature token is absent

from the training vocabulary — a new OS version, browser family, or timezone variant emerging between retraining cycles. `feature_w2v` falls back to the global mean vector for any unseen token, producing a generic embedding that fails to reflect the device’s proximity to the account centroid. `feature_fasttext` computes an embedding for any feature token at inference time by averaging the subword n-gram embeddings for its structured prefix and suffix: `os_harmonyos` shares the `os_` prefix with `os_ios` and `os_android`; `browser_arc` shares the `browser_` prefix with `browser_chrome` and `browser_safari`. The resulting embedding is semantically positioned in the correct region of the feature space, rather than collapsed to the global mean.

Finding 6 — `feature_w2v` improves under Markov sessions (opposite of `id_w2v` in Experiment 2).

In Experiment 2, the ID-based Word2Vec centroid signal degraded under Markov session autocorrelation (AUC 0.982 $\rightarrow$ 0.962), because session runs reduce the effective number of distinct skip-gram training pairs per account, weakening the cluster structure. `feature_w2v` shows the opposite pattern: AUC increases from 0.984 to 0.993 on novel attacks under Markov. The mechanism is distinct. Feature tokens are bounded — approximately 30 distinct values across six dimensions. With stay probability 0.70, the Markov corpus produces many repeated visits from the same device, reinforcing within-account feature co-occurrence patterns. Because the feature vocabulary is small, repeated co-occurrences make account clusters tighter rather than sparser: there is no equivalent of the ID-based device vocabulary growing without bound as accounts accumulate unique device IDs. This divergence is a general result: the corpus structure sensitivity of embedding-based signals depends on whether the token vocabulary is open (device IDs) or bounded (feature tokens). Bounded vocabularies benefit from Markov autocorrelation; open vocabularies are harmed by it.

8. Discussion

8.1 What the three experiments collectively establish

The full set of debate points and their updated verdicts, incorporating Experiment 3:

Point	Question	Exp 1–2 Verdict	Exp 3 Update
§1	OOV detection or account clustering?	Both real, but OOV mechanism inverted	Confirmed: feature-based embeddings demonstrate the genuine clustering mechanism
§2	Do subword n-grams help?	No — they destroy cluster structure (on ID tokens)	Scoped: FastText on structured feature tokens is competitive (AUC 0.985); ruling applies only to opaque ID tokens
§3	Is co-occurrence signal too sparse?	No — negative sampling is sufficient	Confirmed; bounded vocabulary makes it even more robust
§4	Is centroid norm predictive of FP?	No — empirically non-predictive	Unchanged

Point	Question	Exp 1–2 Verdict	Exp 3 Update
§5	Does FastText beat OOV baseline?	No — baseline beats FastText	Reversed context: OOV baseline collapses to 0.750 under corrected evaluation; feature_fasttext achieves 0.985
§6	Threshold on test set problematic?	Minor — ROC-AUC unaffected	Unchanged
§7	ROC-AUC the right metric?	Yes — appropriate for phase one	Unchanged
§8	i.i.d. an upper bound?	Yes — FastText degrades 4× faster	Partially reversed for feature_w2v: bounded vocabulary improves under Markov
§9	Do 64D account clusters exist?	Yes, in Word2Vec. No, in FastText (on ID tokens).	Feature-based embeddings confirm real clusters with stable vocabulary
§10	Enrollment gap a real constraint?	Yes — operational, not theoretical	Reversed: enrollment is a first-class evaluation requirement; OOV binary collapses when enrollment is in the negative class
—	OOV binary as complementary signal	Framed as complementary to Word2Vec centroid	Retired as risk scoring input: evaluation artifact (0.989 AUC) and structural enrollment FP problem (0.750 ceiling)
—	Feature embeddings resolve rotational instability	Proposed in addendum	Confirmed: bounded vocabulary (~30 tokens) is stable across retraining cycles

8.2 Why the original hypothesis survives — with two substitutions

The original PoC used the wrong model and the wrong token type, but tested the right idea. The hypothesis — that per-account device sequences can train embeddings that separate known from novel devices — is validated empirically by Experiments 2 and 3, with two substitutions required.

Substitution 1: Replace FastText with Word2Vec for ID-based embeddings.

FastText’s subword n-gram mechanism, applied to opaque random device ID strings, destroys the block-diagonal co-occurrence structure that skip-gram negative sampling creates. Silhouette goes from -0.051 (FastText, no per-account cluster structure) to $+0.941$ (Word2Vec, near-perfect account separation). AUC goes from 0.910 to 0.982 on i.i.d. data. This substitution is necessary for the fleet-detection signal (`id_w2v`) to work.

Substitution 2: Replace opaque device ID tokens with structured feature tokens, and use FastText on those feature tokens for the primary real-time signal. Using opaque device IDs as tokens means a new device is always OOV, making it indistinguishable from a legitimate enrollment event. Structured feature tokens (`os_ios`, `browser_safari`, `tz_utc-5`) give every device a positioned embedding based on its behavioral fingerprint, not its identifier. When FastText is applied to these structured tokens, its n-gram mechanism reinforces feature-dimension clustering rather than destroying it — because the token prefixes are meaningful and shared across related feature values. This substitution is necessary for the real-time signal to separate enrollment from attack under the corrected evaluation design.

Together, the two substitutions produce the two-path architecture recommended in §9: `feature_fasttext` for real-time profile-fit scoring, `id_w2v` for offline fleet-reuse detection.

8.3 Production constraints that shaped the final architecture

Two production realities were identified during deployment analysis that are outside the scope of the experimental work but materially change the architectural recommendation. Both trace to the same root cause: using opaque device ID strings as embedding tokens.

Rotational instability across training runs. Word2Vec and FastText embedding spaces are not stable across independent training runs. Cosine distance is rotationally invariant within a single model: rotating all vectors by the same orthogonal matrix R leaves all pairwise cosine distances unchanged, because $\text{cos_dist}(Rv_1, Rv_2) = \text{cos_dist}(v_1, v_2)$. The problem is that two independently trained models do not share a coordinate system. Gradient descent from a random initialization finds one of many equivalent local minima; two runs on identical data with different random seeds produce models where the same device receives different vectors — not because its relationships changed, but because the coordinate frame rotated, reflected, and scaled arbitrarily. Formally, if run T produces embeddings V and run $T+1$ produces embeddings V' , there exists some orthogonal matrix R and scale s such that $V' \approx sRV + \epsilon$. An account centroid computed under model version T is a vector in the old coordinate frame; the model it is being compared against at inference time operates in a new, incompatible frame. Every retraining event simultaneously invalidates every account centroid in the system. This has a direct operational requirement: centroids cannot be persisted across model versions. Every retraining event requires a full recompute of every account's centroid.

Note that FastText is not more stable than Word2Vec in this regard. FastText uses a deterministic hash function to map character n-grams to fixed integer indices in the subword embedding matrix — but the learned values stored at those indices are not deterministic. Each training run updates the subword matrix through stochastic gradient descent starting from a different random initialization. The hash table maps the same n-gram string to the same bucket in every run; the value in that bucket changes between runs. FastText and Word2Vec are equally unstable across training runs.

Retraining pressure from legitimate device enrollment. The central production pressure for retraining is device enrollment. A genuine user who buys a new phone has a device that is OOV until the next model retraining cycle incorporates it into the vocabulary. Under any ID-based embedding architecture, that device receives the global mean vector and scores anomalously against the account's centroid, producing false positives from entirely legitimate behavior. At a consumer scale of 10 million accounts, with conservatively 1.5 new device enrollments per account per year across device classes (smartphones replaced every ~ 2.5 years, laptops every ~ 3 – 4 years, tablets every ~ 3 – 5 years), that is approximately 41,000 new device enrollments per day. At a weekly retraining cadence, approximately 280,000 accounts (2.8% of the user base) have a recently enrolled device in the OOV window at any given time, each generating elevated

fraud signals from legitimate behavior. Retraining more frequently reduces the OOV false-positive window but increases the centroid recompute burden and operational complexity — and there is no frequency at which the problem is fully resolved.

Why feature tokens resolve the root cause of both problems. Opaque device ID strings carry no information that transfers across model versions. A device ID like `dev_a3f9...` is assigned its embedding purely by co-occurrence, which depends entirely on the training run’s random initialization. Structured feature tokens carry semantic content: `os_ios`, `browser_safari`, and `tz_utc-5` encode specific behavioral attributes that are stable across model versions. A new device is not OOV under feature-based embedding — it has features, and those features map to existing vocabulary tokens or to well-positioned subword embeddings (via FastText) for new values. The feature vocabulary is bounded (~30 distinct tokens across 6 dimensions), meaning it does not grow as the user base grows or as devices are replaced. Retraining updates the feature-to-embedding mapping without replacing the coordinate frame in a way that invalidates centroids: because the same feature tokens appear in both old and new training runs with similar co-occurrence patterns, the embedding space for feature tokens is more consistent across retraining cycles than the embedding space for opaque device IDs.

Procrustes alignment as mitigation. Mathematically, successive model versions can be aligned by finding the orthogonal rotation matrix $R^* = \operatorname{argmin}_R \|RV' - V\|_F^2$ (subject to $R^T R = I$), solved via SVD as $R^* = UV^T$ where $[U, \Sigma, V^T] = \operatorname{SVD}(V^T V')$. Applying R^* to all new model embeddings maps them into the old model’s coordinate frame, potentially allowing existing centroids to remain valid. This reduces the post-retraining recompute from $O(N \times \text{avg_history_length})$ to $O(\text{anchor_set_size}^2)$ plus incremental updates for newly enrolled devices. However, alignment quality degrades as vocabulary changes grow: for a system adding ~41,000 new devices per day to a vocabulary of 1.8 million tokens (~2.3% daily growth), the anchor set’s shared subspace shrinks with each retraining cycle. Alignment error compounds across retraining cycles — each cycle introduces residual error ϵ_k , and total drift grows faster than linearly for non-stationary data distributions. Architecture changes (embedding dimension, vocabulary tokenization) break alignment entirely. Procrustes alignment is a useful mitigation for low-frequency, stable retraining regimes; it fails for the enrollment pressure described above.

8.4 Limitations

Cross-account device sharing was partially tested in Experiment 3 via fleet injection (80 fleet devices injected into 25% of accounts). The fleet attack AUC of 0.891 for `id_w2v` and 0.920 for `feature_fasttext` represents the first measurement of this confound. Devices shared across accounts for legitimate reasons — family plans, corporate device fleets, VPN endpoints — receive embeddings positioned between multiple account clusters in the ID-based model. Whether this produces acceptable false positive rates at scale, and whether feature-based embeddings are more robust to legitimate sharing, is not yet fully characterized and should be the first targeted stress test.

The returning attacker scenario was not evaluated. An attacker whose device appeared on any account in the training corpus is invisible to any OOV-based signal. Whether `id_w2v` would still flag such a device as foreign to the target account’s cluster depends on the attacker device’s cluster neighborhood — a device co-trained primarily with other accounts’ device sets should land far from the target account’s centroid, but this requires direct empirical measurement. Whether `feature_fasttext` would detect a returning attacker depends on whether the attacker’s feature profile is consistent with the target account’s primary profile.

Realistic class imbalance was not fully characterized. The corrected Experiment 3 evaluation uses a 1:1 ratio of returning known-device events to enrollment events in the negative class, with roughly equal attack volume — not representative of production

class distributions where the positive rate is typically 0.1–1.0%. Precision at fixed recall levels (90%, 95%) under 1,000:1 imbalance is the operationally relevant metric and was not computed. PR-AUC at production fraud rates is the next required measurement.

9. Conclusions and Recommendations

The hypothesis is supported — with two substitutions

The core claim — that per-account device sequences can train embedding models whose centroid distances separate known from novel devices — is empirically validated across three experiments. Two substitutions are required relative to the original PoC.

Substitution 1: For ID-based embeddings, replace FastText with Word2Vec. FastText’s character n-gram mechanism, applied to opaque random device ID strings, destroys the per-account cluster structure (silhouette -0.051) that the approach depends on. Word2Vec preserves it cleanly (silhouette $+0.941$, AUC 0.982 on i.i.d., AUC 0.962 on Markov). The ID-based Word2Vec signal (`id_w2v`) achieves AUC 0.891 on fleet/reuse attacks and is the correct offline batch signal for that threat model.

Substitution 2: For the primary real-time signal, replace opaque device ID tokens with structured feature tokens and apply FastText to those tokens (`feature_fasttext`). Using device IDs as tokens makes any new device OOV, producing false positives on legitimate enrollment that are structurally indistinguishable from attacks — the OOV binary baseline’s 0.989 AUC in Experiment 2 was entirely an artifact of the evaluation design, not a real discriminative signal. Under the corrected evaluation (enrollment in the negative class), all OOV binary signals collapse to an analytically determined ceiling of AUC 0.750 on novel and spoof attacks, and are anti-correlated with fleet detection (AUC 0.250). `feature_fasttext` achieves AUC 0.985 on novel attacks and AUC 0.920 on fleet attacks with enrollment in the negative class, because it measures profile fit rather than device ID novelty.

Recommended architecture

The production architecture operates on two paths at different timescales, with a separate operational gate.

Real-time signal — `feature_fasttext` (feature-profile proximity):

The primary real-time signal is cosine distance between a device’s feature-based FastText embedding and the account’s feature centroid. Each login event maps to a sequence of six structured feature tokens (OS, browser, timezone, language, network type, screen resolution). The signal measures whether the incoming device’s behavioral fingerprint fits the account’s established profile — not whether the device ID has been seen before.

This signal is recommended for production for four reasons. First, it structurally separates legitimate enrollment from attack: enrollment devices have profiles consistent with the account’s primary profile and land close to the account centroid; attack devices with foreign profiles land far. The OOV binary baseline cannot make this distinction at any threshold. Second, it handles OOV feature tokens at inference time via n-gram averaging over structured token prefixes (`os_`, `browser_`, `tz_`): when a new OS version, browser family, or timezone variant appears between retraining cycles, `feature_fasttext` produces a positioned embedding for it rather than a global mean fallback. Third, the bounded feature vocabulary (~30 distinct tokens across 6 dimensions) is rotationally stable across retraining cycles: the same feature tokens appear in successive training runs with similar co-occurrence patterns, so centroids computed in one model version remain

approximately valid in the next — unlike ID-based embeddings, where every retraining event invalidates all stored centroids. Fourth, AUC is 0.985 on novel attacks, 0.920 on fleet-reuse attacks, and 0.798 on spoof attacks with enrollment in the negative class — the only single signal that achieves strong performance across all three attack types simultaneously.

Offline batch signal — `id_w2v` (ID-based Word2Vec centroid distance):

Word2Vec skip-gram trained on per-account device ID histories, with cross-account fleet device injection, achieves AUC 0.891 (i.i.d.) / 0.856 (Markov) on fleet/reuse attacks. Fleet devices trained across multiple accounts have embeddings positioned in a cross-account region of the space, not close to any single account’s centroid. When a fleet device appears against a new target account, the cosine distance to that account’s centroid is high. This signal targets a real and distinct threat model — device reuse across victims — that ID-novelty-based signals cannot detect.

This signal is not suitable for real-time inference. ID-based embedding spaces are not stable across training runs: every retraining event requires a full recompute of all account centroids. The correct deployment is an offline batch pipeline that recomputes centroids after each retraining cycle and scores accounts daily or weekly, routing elevated-distance patterns to a review queue. The signal is also not suitable as the sole detection mechanism: `id_w2v` scores OOV devices (novel attacks, enrollment events) identically at approximately the global mean’s distance from the account centroid, making it structurally unable to distinguish novel attacks from legitimate enrollment. The two signals address complementary threat models and should be deployed on separate timescales.

Operational gate — per-account known-device set:

A per-account known-device set stored in a low-latency database (Redis, DynamoDB) serves as the operational gate for step-up authentication. It is updated in milliseconds when a device is confirmed legitimate (step-up authentication completed, device management enrollment) or confirmed fraudulent (fraud review). It is decoupled from model training entirely: no retraining is required to enroll or revoke a device. This is the correct trigger for step-up authentication — not a risk score input, but an operational flag that prompts further verification. It must not be used as a risk scoring signal: its score encodes “device ID novelty,” not “behavioral fit to this account,” and collapses to the 0.750 ceiling under the corrected evaluation.

Implementation priorities

Priority	Action	Rationale
1	Implement <code>feature_fasttext</code> as the primary real-time signal	Structurally separates enrollment from attack; handles OOV feature tokens at inference; bounded vocabulary is stable across retraining; AUC 0.985 on novel attacks, 0.920 on fleet reuse with enrollment in negative class
2	Replace FastText with Word2Vec for ID-based embedding (<code>id_w2v</code>)	7-point AUC gain for fleet/reuse detection (0.910 → 0.982); silhouette 0.941 vs -0.051; treat as offline batch signal, not real-time

Priority	Action	Rationale
3	Move ID-based centroid scoring to an offline batch pipeline	Embedding spaces are not stable across retraining runs; centroid recompute cannot be done in real time at scale
4	Retain per-account known-device set as an operational gate	Step-up auth trigger on OOV events; decoupled from risk scoring; millisecond update latency on enrollment confirmation
5	Stress-test cross-account device sharing	Highest-impact unresolved confound for both signals; fleet injection at 25% of accounts is a starting point but does not cover legitimate sharing scenarios
6	Evaluate the returning-attacker scenario	Critical gap; id_w2v offline signal is primary coverage for this case; requires direct empirical measurement
7	Report PR-AUC at realistic imbalance ratios	ROC-AUC validated the signal ranking; operational cost requires PR-AUC at 0.1–1.0% positive rate
8	Evaluate feature_w2v + id_w2v ensemble	Combined signal should outperform either alone; logistic regression over both targets complementary threat models
9	Extend spoof attack coverage with higher-cardinality features	feature_fasttext AUC 0.798 on spoof (timezone mismatch only); test harder spoof with more matched dimensions; quantify the minimum feature mismatch needed for detection

Appendix — Artifacts

File	Contents
ato_fasttext_poc.py	Original PoC (FastText, i.i.d. corpus)
ato_experiment2.py	Experiment 2 (FastText, Word2Vec, OOV baseline; i.i.d. + Markov)
ato_experiment3.py	Experiment 3 (fleet corpus, feature embeddings, corrected enrollment evaluation)
plot_conclusions.py	Generates Experiment 2 figures
CRITIQUE.md	Ten-point systematic critique
DEFENSE.md	Point-by-point rebuttal
DEBATE.md	Multi-turn debate to resolution per point

File	Contents
CONCLUSIONS.md	Findings from Experiments 2 and 3, with debate scorecard and signal hierarchy
REPORT_ADDENDUM.md	Production deployment analysis: rotational instability, retraining pressure, Procrustes alignment, architecture
README.md	Quickstart and caveats

All scripts are self-contained and runnable with `uv run <script>.py`.

All experiments use $SEED = 42$. Reported numbers are stable across runs to ± 0.005 AUC.